



Iroko

Predictive Datacenter Congestion Control

Congestion Control in Data Centers

- Ideal data center should have:
 - Low latency.
 - High utilization.
 - No packet loss or queuing delay.
 - Fairness.
- Most congestion control and flow scheduling variations are based on TCP.
 - Prioritizes **fairness** and **utilization**.
 - End hosts receive feedback based on round-trip latency and packet loss.
 - Hosts gradually adjust to “**fair**” network rate.

Problem

- Data centers are continuously improving
- TCP is a fundamentally **reactive** protocol.
 - When TCP's backoff kicks in, network conditions are already suboptimal.
 - TCP's convergence rate for small and medium sized flows can not keep up.
- A reactive protocol can not prevent queue buildup and bursts.
 - **Queueing latency** dominates.
 - Frequent retransmits reduce **goodput**.
 - Data center performance may be **unstable**.

Datacenter Network

Small Latency

$< 100 \mu s$



Shallow Buffer

$< 30 \text{ MB}$ for ToR




High Bandwidth

10/40 ~ 100 Gbps



Large Scale

$> 10,000$ machines



Current research insights

- Admission control can have value over burst-and-backoff.
 - **ExpressPass** (2017): Network-enforced credit packet protocol.
 - **FastPass** (2014): “Zero-Queue” network managed by a central arbiter.
- Centralized traffic control solutions may be able to **scale** sufficiently.
 - Datacenter controllers have sufficient computing power at disposal (FastPass, 2014).
 - Benefit of **global traffic knowledge** amortizes RTT cost (Hedera, 2010).
 - **Compartmentalization** enables localized traffic management (Jupiter, 2015).
- **Preemptive** scheduling is possible in data centers.
 - Certain traffic patterns are **predictable** (MicroTE, 2011; Data Centers in the Wild, 2010).
 - Congestion resource management can be **trained** (Remy, 2014).

Our Idea

- What if we build an “**intelligent**” traffic arbiter?
 - **We** control the sending rate of hosts.
 - Arbiter **decides** bandwidth allocation per host.
 - Continuously **learns** on traffic feedback.
 - Tries to **predict** sending behaviour.
 - Schedules **future** host tx rate.
- Three research questions:
 - How would such a controller look like?
 - Is it really possible to analyze and predict data center traffic?
 - How does this design compare to more conservative solutions?



A traffic controller.


Iroko



Architecture

receive
information




Send packet to update
bandwidth



Implementation: Topology Simulation

- Mininet:
 - Instant Virtual Network
- Fat Tree Topology:
 - 20 switches
 - 16 hosts
- Traffic Simulation:
 - Use Hedera benchmarks

> sudo mn



Implementation: Data Acquisition

- Decide to use data being collected from the connections between interfaces
 - Interfaces known in advance
- Use linux commands to get information directly:
 - tc qdisk & awk
- Allows us to collect usable information of network:
 - bandwidth, free bandwidth, drops, overlimits, and queues

```
cmd = "awk \"/^ *%s: / '\"' { if ($1 ~ /.*: [0-9][0-9]*/) { sub(/^.*:/, '\"') ; print $1 } else { print $2 } }\"' /proc/net/dev" % (
    iface)
try:
    output = subprocess.check_output(cmd, shell=True)
except:
    print("Empty Request")
    output = 0
```




Implementation: Communication with hosts

- Arbiter needs to be able to control hosts bandwidth
 - Otherwise it isn't all powerful
- Logic is implemented in application layer
 - C++ communicating via python sockets
- High Level View:
 - Controller sends packet \Rightarrow packet tells host how much traffic can send out



Implementation: Learning Model

- Prototype idea is simple temporal difference update:
 - $V(s) = V(s) + A_{\text{learningRate}} * (R + V(s') - V(s))$
 - R = modifier based on loss
- Goal Implementation:
 - Actor-Critic Model
 - Use loss rate & utilization as indication of success




Evaluation & Methodology

- Metrics we are mainly interested in
 - Utilization
 - Packet drop rate
 - Queuing delay
 - Fairness
 - Starvation
- Target environment
 - Mininet simulation of small data center of 16 hosts and 20 switches with FatTree topology
 - ECMP enabled switches for load balancing
 - Assume switches support OpenFlow
- Comparison with Hedera, DCTCP and ECMP
 - TCP Cubic for Hedera and ECMP
 - Using UDP traffic for Iroko. Rate limiting implemented at the application layer

Iroko vs ECMP vs Hedera vs DCTCP


- ECMP
 - Switch hashes the <src ip, src port, dst ip, dst port, protocol> to a port
 - Packets belongs to same flow routes to same port
- Hedera
 - Central controller pulls traffic information from the switches
 - Identifies “elephant flows” and assigns dedicated path
- DCTCP
 - TCP optimized for data center
 - Congestion window size back off based on ECN
 - Marks packets passing through a link with utilization more than a threshold

Preliminary Results (Utilization)




Discussion

- ML is a tricky business:
 - What **learning model** to choose?
 - What **parameters** to choose?
 - **Convergence takes time**; need to learn for a long time to see effects that repeat over long time scales
 - Are we chasing our own tail?
- Is datacenter traffic actually **predictable**?
 - Elephant flows might not need to be predicted. Hosts can report.
 - Short lived flows might be very hard to predict and may need finer grained epochs.
 - Is datacenter traffic ergodic (statistical properties can be deduced from a sufficiently long random sample) ?



Discussion Continued



- Scaling?
 - **Simulation network is small**
 - How does learning scale?
 - Learning space is huge
- Eval Limitations:
 - What is the deciding factor in the performance (many variables)
 - Simulation is on local machine; **OS scheduling and utilization effects** come into play
 - Reproducibility and correctness of the implementation
- Topology and Traffic Matrix
 - Currently using a (small) fat-tree topology
 - **Predetermined traffic matrices.** May not be representative of real traffic
 - How do the techniques translate to other topologies and different traffic matrices

Future Work

- Add explicit **host feedback**
 - Hosts can actively request more data
 - Hosts can send **predictive requests** e.g. Map-Reduce job is expected to finish shortly
- Learning Improvements
 - Learn on $t \cdot \text{epoch_length}$; predict and update hosts on $(1-t) \cdot \text{epoch_length}$ where t in $(0,1)$
 - Learn on individual flows rather than hosts
- Use a **token based system** to allocate bandwidth
 - Allows hosts to make choices about when to use bandwidth
 - Can predict hosts anticipated future bandwidth needs by tokens kept in reserve by hosts
 - Host can cooperate - use techniques from game theory to organize
- Different routing strategy and traffic matrices
 - Currently relying on ECMP for routing; explore other routing options
 - Acquire actual datacenter traffic to test on