

TCPTokens: Introducing Currency into Data Center Congestion Control

Anand Jayarajan

University of British Columbia
anandj@cs.ubc.ca

Robert Reiss

University of British Columbia
rreiss@cs.ubc.ca

Michael Przystupa

University of British Columbia
michael.przystupa@gmail.com

Fabian Ruffy

University of British Columbia
fruffy@cs.ubc.ca

ABSTRACT

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Vestibulum lacinia arcu ac diam mattis pharetra. Curabitur luctus posuere dictum. In quam tortor, cursus luctus tortor ut, porta fermentum lacus. Nunc eget augue dui. Proin ante lectus, malesuada a sem ut, lobortis eleifend turpis. Morbi sit amet odio interdum, sollicitudin sem eget, mollis orci. Curabitur rhoncus arcu magna, eget rhoncus nunc porttitor at. Nulla quis posuere urna. Nam enim nunc, hendrerit et gravida et, mollis a urna. Donec aliquam massa sollicitudin tortor cursus placerat. Ut non nisi purus. Etiam laoreet ut libero et tempus.

KEYWORDS

TCP, Congestion Control, SDN, Data center

1 INTRODUCTION

In the early years, research in network congestion control has been dominated by the traditional assumption of a decentralised, autonomous network. End-hosts only have control over the amount of traffic they send, and are unaware of the intentions or traffic rate of their peers. Similarly, switches and routers are unaware of global traffic patterns and only forward based on their local notion of optimality.

In line of these assumptions, TCP has been designed to optimize traffic globally on a simple fairness principle. Nodes react to packet loss and assume that others will do the same. An equilibrium is reached when all end-hosts achieved a traffic rate that, in sum, conforms to the maximum available bandwidth of the congested link.

TCP works excellently in scenarios where many different distrustful participants compete for limited bandwidth. Still, TCP is a reactive protocol. The fact that packet loss and latency increase occur in the network, already indicates a problem. Packet loss is largely caused by an overflow of queues ("bufferbloat") in forwarding elements, implying that traffic has not been optimally distributed.

Ideally, a network should always be "zero-queue", i.e., latency will merely be caused by propagation, and not queuing delay. Queueing has generally not been a dominant issue in wide-area and enterprise networks, as traffic is sufficiently distributed and diverse, with only few "hot" target hosts. Optimal traffic optimization is a substantial challenge, if not impossible, as network operators have no control

over the individual network elements and its participants. Under these conditions TCP and its extensions can be considered a best-effort solution.

However, these assumptions have changed in the past decade due to new developments. Datacenters have emerged as a driving force in networking research, as they pose new design challenges.

They are driven by minimization of costs and maximization of compute power and are intended to run at maximum utilization to achieve an optimal compute/cost ratio. In such a scenario, inefficient routing can lead to a quick deterioration and collapse of a high-load network, requiring more sophisticated approaches to solve congestion control. However, operators now also have the ability to freely control and adapt their network architecture, leading to highly customized systems and fine-grained control.

Much in line with the trend of datacenters, Software-Defined Networking (SDN) emerged as a new networking paradigm. Moving away from the principle of distributed communication and routing, SDN introduced the notion of "centralized management". A single controller with global knowledge is able to automatically modify and adapt the forwarding tables of all switches in the network, as well as notify end hosts of changes in the network. These two new trends in systems allowed for impactful new innovation in the space of TCP congestion research. Traffic can now be managed centralised based on global knowledge of the entire topology and traffic patterns. In recent years, a new line of centralised schedulers has emerged, which make use of these advantages and achieve close to optimal bandwidth optimization.

However, these schedulers are fundamentally still reactive. The central controller responds to changes in the network or requests by applications, which may cost valuable roundtrip latency. Often, short-term flows or bursts are unaccounted for which causes undesirable packet loss and back-propagating congestion.

A much more desirable solution is a global, centralised arbiter which is able to predict and fairly distribute flows in the network before bursts or congestion occurs. By treating the network's compute and forwarding power as a single finite resource, a controller could act like the OS scheduler distributing CPU time slices to processes. This design approach also follows SDN's aspiration of introducing operating systems abstractions to the networking domain space.

In this project, we plan to explore the possibilities of a centralised, proactive flow scheduler. We ask ourselves the following research questions:

- (1) Is it possible to design a centralised token-based scheduling network?
- (2) Is it possible to predict traffic and pre-emptively schedule flows and token distribution in a datacenter context?
- (3) Using this approach, are we able to achieve better performance and utilization than existing solutions?

In the scope of this course, we attempt to answer question 1 and design a simple token-based scheduler in Mininet. If we succeed, we will benchmark our results and evaluate the level of utilization compared to contemporary scheduling systems.

2 RELATED WORK

Nulla facilisi. Donec congue, quam sed faucibus varius, libero dolor mollis ipsum, vel placerat felis urna quis nibh. Morbi tempor ipsum nec urna porttitor interdum. Phasellus cursus nisi nec tempor blandit. Curabitur at sem convallis, sagittis purus ornare, sagittis urna. Nullam eget eleifend orci. Suspendisse vel justo at elit mollis porta. Nulla congue sapien nibh, tincidunt pulvinar orci elementum in. Aliquam mollis tempus velit non pretium. Aliquam hendrerit at orci id bibendum. Praesent gravida sapien eu eleifend faucibus.

Sed at tristique augue. Proin non venenatis lacus. Sed sodales ornare urna, id vulputate dui dictum in. Duis blandit, ex in ornare vulputate, est nibh ornare felis, vitae egestas libero leo nec enim. Fusce non arcu id eros suscipit feugiat. Nullam dapibus nibh a vestibulum maximus. Fusce convallis tortor odio, vitae egestas quam lacinia vitae. Aliquam laoreet convallis eros vel ultricies. Morbi non orci nunc. Proin accumsan posuere felis vitae pharetra. Aenean suscipit congue felis, in sollicitudin metus aliquam et. Pellentesque at commodo mauris. Morbi aliquet venenatis nisl, a cursus ligula cursus et. Cras sit amet lorem ac turpis vestibulum ultrices. Praesent cursus nibh vel malesuada convallis.

3 IMPLEMENTATION

Nunc ornare tortor at finibus faucibus. Sed bibendum efficitur nibh ac ultricies. Donec maximus purus lectus, id auctor diam molestie vel. Pellentesque ex odio, ultricies vel magna et, laoreet placerat augue. Suspendisse nunc ipsum, semper rhoncus sem at, dignissim mollis mauris. Cras viverra, eros eget faucibus porttitor, quam ex aliquet mi, eget sodales elit dolor in mi. Aenean congue augue id nunc blandit, ac molestie augue sagittis. Duis rhoncus odio eu arcu posuere, eget luctus diam lobortis. Aliquam sed ullamcorper velit. Suspendisse sed ligula sed risus rhoncus placerat ut fermentum nulla. Nam facilisis dui metus, non malesuada lacus feugiat quis. Vivamus volutpat non nibh eu gravida. Suspendisse efficitur leo maximus tortor interdum maximus. Nulla feugiat elit nec tincidunt tincidunt. In fermentum lacus tellus, sit amet tempor leo pulvinar quis.

Duis tortor mauris, dignissim vehicula ligula non, lobortis finibus odio. Curabitur eu lobortis erat. Donec arcu ex, cursus in finibus sit amet, commodo ac mi. Integer mattis iaculis ante placerat tempus. Phasellus eu diam blandit, ultricies mauris ut, viverra odio. Praesent accumsan, lectus ut accumsan semper, sem elit pharetra enim, at vehicula felis ante a tortor. Duis at tellus id enim mollis posuere et vitae nisi. Nam convallis vulputate commodo. Vestibulum eget libero sed eros dictum facilisis. Praesent posuere vulputate ipsum non sagittis.

4 EVALUATION

Maecenas aliquet sagittis nunc, ac sagittis justo scelerisque in. Maecenas at libero leo. Integer maximus gravida ante, accumsan tempor dui. Quisque accumsan semper erat, at mattis ex vestibulum at. Proin eu nulla vitae ex molestie facilisis eu id dolor. Phasellus lobortis efficitur purus, sit amet laoreet lacus sodales vulputate. Maecenas venenatis est a justo blandit consequat. Donec vitae pharetra orci. Duis tortor diam, interdum sed ex non, ultricies suscipit lorem. Pellentesque eros arcu, dignissim ut leo sit amet, sodales sodales tortor. Ut convallis, elit nec auctor mollis, erat neque feugiat ex, bibendum pharetra nulla risus sed neque. Aliquam feugiat elit metus, non convallis nisl malesuada pharetra. Donec ligula neque, blandit eget ligula placerat, lacinia finibus odio. Quisque sit amet suscipit diam.

Nunc consectetur turpis non sollicitudin pellentesque. Nullam nec eleifend lectus. Nunc sem enim, bibendum quis elit nec, laoreet posuere dui. Vivamus sodales nisi sed facilisis consectetur. Donec venenatis justo sed consectetur ullamcorper. Nulla dignissim magna eget arcu consectetur aliquet. Praesent consequat, elit quis auctor egestas, dolor mi pretium erat, eu ullamcorper risus velit sed metus. Praesent lectus urna, tempor sed imperdiet id, condimentum id nunc. Aliquam id mi sollicitudin, venenatis justo id, congue diam. Sed quis dictum odio. Donec auctor enim quis ultricies consequat. Cras maximus a velit luctus feugiat.